Peer Response 1

Hi Gesine,

Your analysis of the ACM case study provides an excellent overview of the ethical and legal dimensions of Max and Jean's conduct. I particularly appreciated your integration of both ACM and BCS codes, which effectively demonstrates how professional standards converge on fundamental principles of dignity and respect.

I would like to expand on your point regarding Jean's failure to act. This illustrates the concept of ethical passivity, where inaction becomes complicity in perpetuating harmful workplace cultures. Shahzad and Sandhu (2025) argue that organisational ethical climates are fundamentally shaped by leadership behaviours and the interventions they make, or fail to make, in addressing misconduct. Jean's dismissal of Diane's concerns signals to the entire team that such behaviour is tolerated, actively undermining trust and psychological safety within the digital workplace.

Furthermore, the gendered dimension you identified is particularly concerning given persistent underrepresentation of women in computing roles. Kuteesa *et al*. (2024) found that 50% of women in STEM jobs reported experiencing gender discrimination, with workplace cultures often characterised by hostile environments that drive talented women from the industry. Allowing Max's behaviour to continue unchallenged not only harms individuals but actively contributes to the systemic barriers preventing diversity in technology professions, ultimately limiting innovation and organisational growth.

Word count: 206

References:

Kuteesa, K.N., Dagunduro, A.O., Ajuwon, O.A. and Chikwe, C.F. (2024) 'Addressing gender inequality in tech workplaces: challenges and opportunities', *Comprehensive Research and Reviews in Multidisciplinary Studies*, 2(1), pp. 19-26. Available at: https://doi.org/10.57219/crrms.2024.2.1.0037

Shahzad, M.U. and Sandhu, M.R.S. (2025) 'Organizational ethical climate: a systematic review of two decades exploring trust, leadership, and sustainable practices in the digital era', *Lahore Business & Economics Journal*, 4(2), pp. 111-136. Available at: https://pu.edu.pk/home/journal/52/v4_2_2025.html (Accessed 29 October 2025).

Peer Response 2

Hi Nelson,

Your comprehensive analysis effectively highlights the ethical failures in AllTogether's development process. The leadership's dismissive question "How many blind people use our product, anyway?" particularly exemplifies the problematic mindset that continues to plague software development (Horton, 2024).

However, the case reveals a deeper systemic issue beyond individual ethical lapses. Recent research by da Silveira *et al.* (2024) identifies that organisations frequently lack strategic motivations and structured commitments to integrate accessibility throughout development lifecycles, even when developers possess technical knowledge. This suggests that AllTogether's failure stemmed not merely from scheduling pressures but from absent organisational frameworks that would prioritise accessibility as a core requirement rather than an optional feature.

The knowledge gap you mentioned deserves further emphasis. Studies demonstrate that computing education inadequately prepares professionals to address accessibility challenges, creating substantial skill deficits across the industry (El Morr *et al.*, 2024). This educational shortcoming explains why AllTogether's product team failed to recognise accessibility violations during development, they likely lacked the expertise to identify issues before release. Consequently, organisations increasingly rely on dedicated accessibility specialists, yet digital inaccessibility persists despite growing professional certification and expertise availability.

Your conclusion regarding proactive embedding of accessibility is crucial. The case demonstrates that reactive approaches generate cascading costs: emergency remediation, customer dissatisfaction, legal exposure, and damaged professional credibility.

Word count: 214

References:

Horton, S. (2024) *Case Study: Accessibility in Software Development*. Available at: https://www.acm.org/code-of-ethics/case-studies/accessibility-in-software-development (Accessed: 29 October 2025).

da Silveira, L.A., Aljeedani, W. and Eler, M.M. (2024) 'Strategic Insights into Integrating Accessibility in Software Development: Motivations, Barriers, Commitments, and Business Value', *Proceedings of the XXIII Brazilian Symposium on Human Factors in Computing Systems*, Brasilia DF, Brazil, 7-11 October. New York: Association for Computing Machinery. Pp. 1–14. Available at: https://doi.org/10.1145/3702038.3702108

El Morr, C., Singh, D., Sawhney, V., Fernandes, S., El-Lahib, Y. and Gorman, R. (2024) 'Exploring the intersection of AI and inclusive design for people with disabilities', *Studies in Health Technology and Informatics*, 316, pp. 556-559. Available at: https://doi.org/10.3233/shti240475

Peer Response 3

Hi Nikolaos,

Thank you for your analysis of the Blocker Plus case. It effectively demonstrates how machine learning systems remain vulnerable to adversarial manipulation. The activist groups' exploitation of the feedback mechanism represents a classic data poisoning attack, where malicious actors deliberately inject biased training data to corrupt model behaviour.

This vulnerability reveals a fundamental challenge in contemporary content moderation systems. Research by Gomez *et al.* (2024) demonstrates that even state-of-the-art toxicity detection models exhibit substantial arbitrariness, with different implementations producing inconsistent classifications on identical content. This predictive multiplicity creates opportunities for manipulation whilst simultaneously generating disparate impacts across social groups, precisely what occurred when Blocker Plus began censoring legitimate content on LGBTQ+ rights and vaccination. The developers' failure to implement robust validation mechanisms exemplifies how technical oversight translates directly into discriminatory outcomes.

Your emphasis on transparency and accountability aligns with emerging governance frameworks. Recent systematic reviews identify these principles as central to responsible AI deployment, alongside human oversight and explainability (Ribeiro *et al.*, 2025). However, Robles and Mallinson's (2025) unified governance framework suggests that fragmented ethical guidelines prove insufficient without enforceable accountability structures. The Blocker Plus developers possessed ethical codes yet failed to act. This highlights that professionalism requires proactive risk assessment, adversarial robustness testing, and transparent incident response protocols embedded within organisational practice, not merely aspirational principles.

Word count: 219

References:

Gomez, J.F., Shukla, A., Gómez, S. and Ruiz, C. (2024) 'Algorithmic arbitrariness in content moderation', *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency (FAccT '24)*, Rio de Janeiro, Brazil, 3-6 June .New York: Association for Computing Machinery. Pp. 2234-2253. Available at: https://doi.org/10.1145/3630106.3659036

Ribeiro, D., Rocha, T., Pinto, G., Cartaxo, B., Amaral, M., Davila, N. and Camargo, A. (2025) 'Toward effective AI governance: a review of principles', *arXiv*. Available at: https://doi.org/10.48550/arXiv.2505.23417

Robles, P. and Mallinson, D.J. (2025) 'Advancing AI governance with a unified theoretical framework: a systematic review', *Perspectives on Public Management and Governance, gvaf013*, pp. 1-15. Available at: https://doi.org/10.1093/ppmgov/gvaf013.